

# Experiment 3 : EDA on *income-dataset.csv*

## load dataset

```
In [12]: import pandas as pd
column_names = [
    "age", "workclass", "fnlwgt", "education", "education_num",
    "marital_status", "occupation", "relationship", "race", "gender",
    "capital_gain", "capital_loss", "hours_per_week", "native_country", "income"
]
df = pd.read_csv('income-dataset.csv', header=None, names = column_names)
df.head()
```

Out[12]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relat
<b>0</b>	39	State-gov	77516	Bachelors		13	Never-married	Adm-clerical
<b>1</b>	50	Self-emp-not-inc	83311	Bachelors		13	Married-civ-spouse	Exec-managerial
<b>2</b>	38	Private	215646	HS-grad		9	Divorced	Handlers-cleaners
<b>3</b>	53	Private	234721		11th	7	Married-civ-spouse	Handlers-cleaners
<b>4</b>	28	Private	338409	Bachelors		13	Married-civ-spouse	Prof-specialty

1. What are the key attributes in the given dataset, and what types of data do they contain?

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
  0   age              32561 non-null   int64  
  1   workclass        32561 non-null   object  
  2   fnlwgt           32561 non-null   int64  
  3   education        32561 non-null   object  
  4   education_num    32561 non-null   int64  
  5   marital_status   32561 non-null   object  
  6   occupation       32561 non-null   object  
  7   relationship     32561 non-null   object  
  8   race              32561 non-null   object  
  9   gender            32561 non-null   object  
  10  capital_gain    32561 non-null   int64  
  11  capital_loss    32561 non-null   int64  
  12  hours_per_week  32561 non-null   int64  
  13  native_country   32561 non-null   object  
  14  income            32561 non-null   object  
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

In [14]: `df.describe()`

Out[14]:

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_
<b>count</b>	32561.000000	3.256100e+04	32561.000000	32561.000000	32561.000000	3256
<b>mean</b>	38.581647	1.897784e+05	10.080679	1077.648844	87.303830	4
<b>std</b>	13.640433	1.055500e+05	2.572720	7385.292085	402.960219	7
<b>min</b>	17.000000	1.228500e+04	1.000000	0.000000	0.000000	1
<b>25%</b>	28.000000	1.178270e+05	9.000000	0.000000	0.000000	2
<b>50%</b>	37.000000	1.783560e+05	10.000000	0.000000	0.000000	4
<b>75%</b>	48.000000	2.370510e+05	12.000000	0.000000	0.000000	5
<b>max</b>	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	9

In [15]: `# Check categorical features`  
`print(df.describe(include='object'))`

	workclass	education	marital_status	occupation	relationship	\
<b>count</b>	32561	32561	32561	32561	32561	32561
<b>unique</b>	9	16	7	15	6	
<b>top</b>	Private	HS-grad	Married-civ-spouse	Prof-specialty	Husband	
<b>freq</b>	22696	10501	14976	4140	13193	

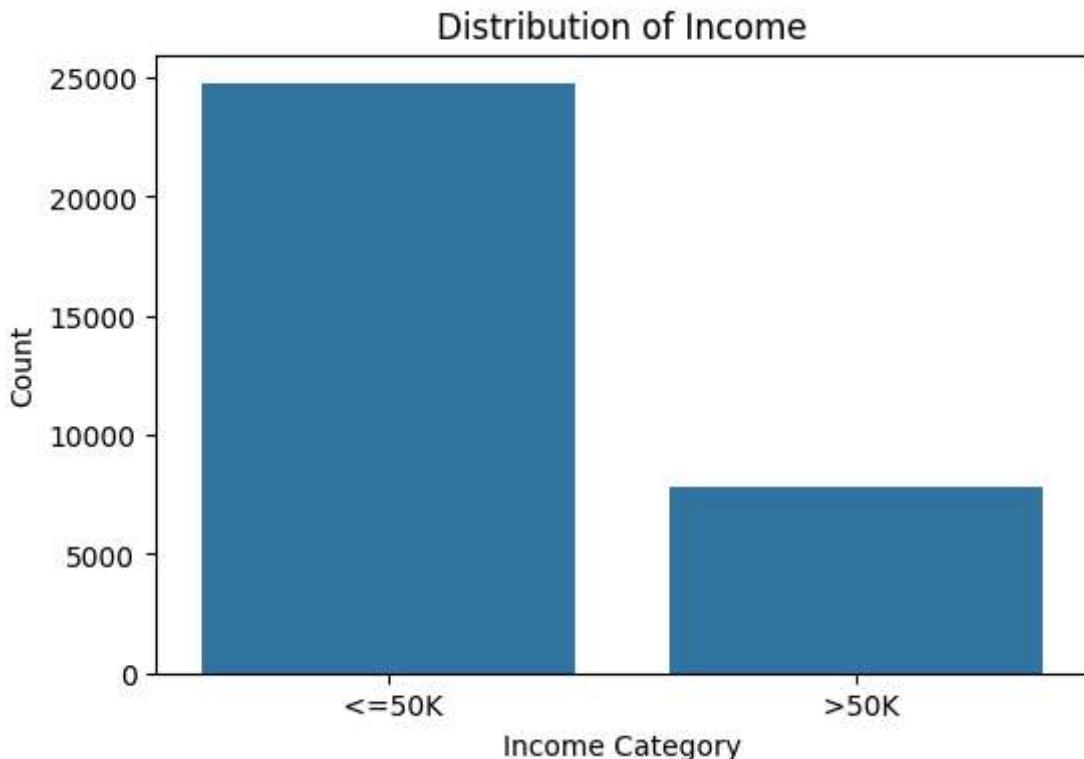
  

	race	gender	native_country	income
<b>count</b>	32561	32561	32561	32561
<b>unique</b>	5	2	42	2
<b>top</b>	White	Male	United-States	<=50K
<b>freq</b>	27816	21790	29170	24720

## 2. What is the distribution of income in the dataset? (Use histograms or boxplots for visualization)

```
In [17]: import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(6,4))
sns.countplot(x='income', data=df)

plt.title("Distribution of Income")
plt.xlabel("Income Category")
plt.ylabel("Count")
plt.show()
```



## 3. Are there any missing or null values in the dataset? How would you handle them?

```
In [18]: df.isnull().sum()
```

```
Out[18]: age          0
          workclass      0
          fnlwgt         0
          education       0
          education_num   0
          marital_status  0
          occupation      0
          relationship    0
          race            0
          gender           0
          capital_gain    0
          capital_loss    0
          hours_per_week   0
          native_country   0
          income           0
          dtype: int64
```

there's no null values in our dataset.

Missing values in a dataset can be handled in several ways. You can remove the rows or columns if the number of missing values is small. For numerical data, missing values can be replaced with the mean or median. For categorical data, they can be filled with the most frequent value (mode). Sometimes, missing values can be replaced with a constant such as "Unknown." In advanced cases, predictive methods like KNN imputation can be used. The method chosen depends on the amount of missing data and the type of variable.

#### 4. What is the correlation between different numerical attributes? Use a heatmap to visualize this.

```
In [26]: import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(15,12))
sns.heatmap(df.select_dtypes(include=['Int64']).corr(), annot=True)
plt.title("Correlation Heatmap of Numerical Attributes")
plt.show()
```



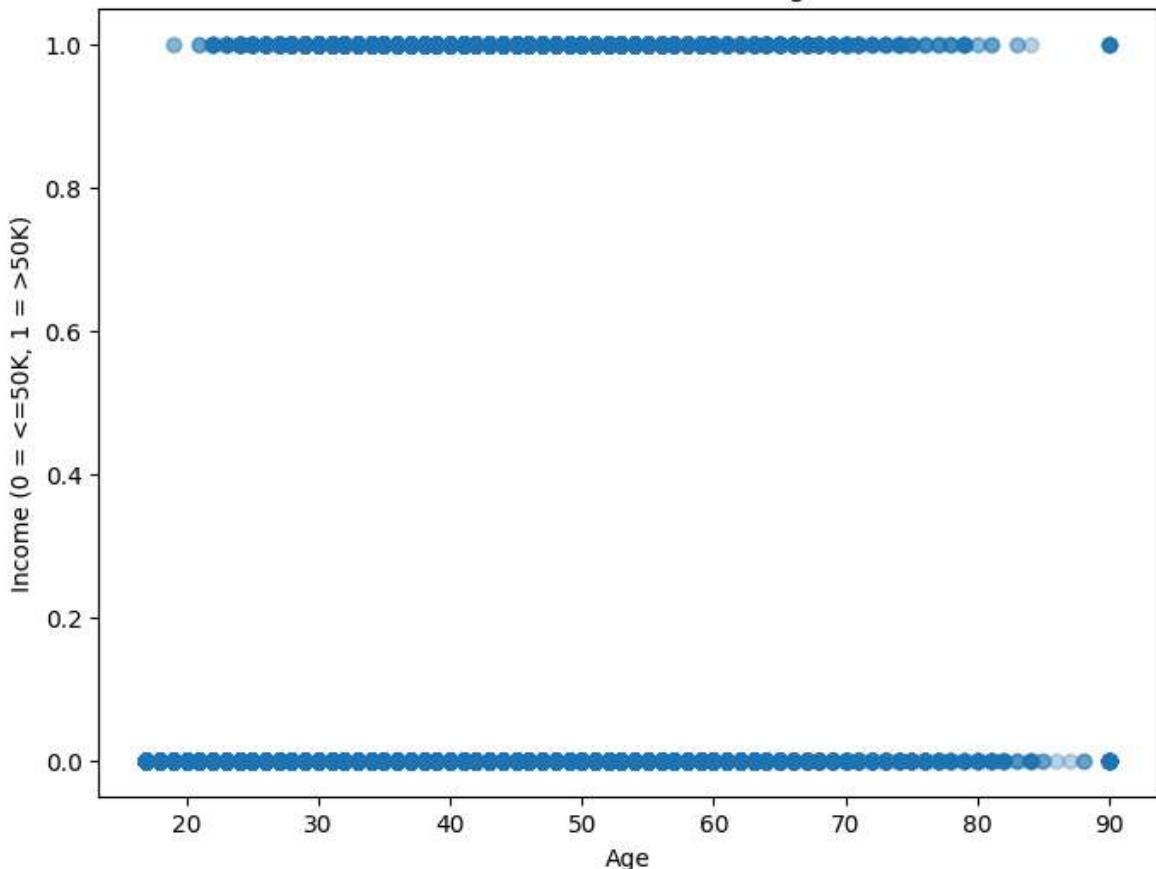
## 5. How does income vary with respect to age? Generate a scatter plot to analyze the relationship.

```
In [20]: df['income'] = df['income'].str.strip()

# Convert income to binary
df['income_binary'] = df['income'].apply(lambda x: 1 if x == '>50K' else 0)

# Scatter plot
plt.figure(figsize=(8,6))
plt.scatter(df['age'], df['income_binary'], alpha=0.3)
plt.xlabel("Age")
plt.ylabel("Income (0 = <=50K, 1 = >50K)")
plt.title("Scatter Plot: Income vs Age")
plt.show()
```

Scatter Plot: Income vs Age



## 6. Which profession has the highest and lowest average income? Visualize this using a bar chart.

```
In [21]: # Clean whitespace
df['income'] = df['income'].str.strip()

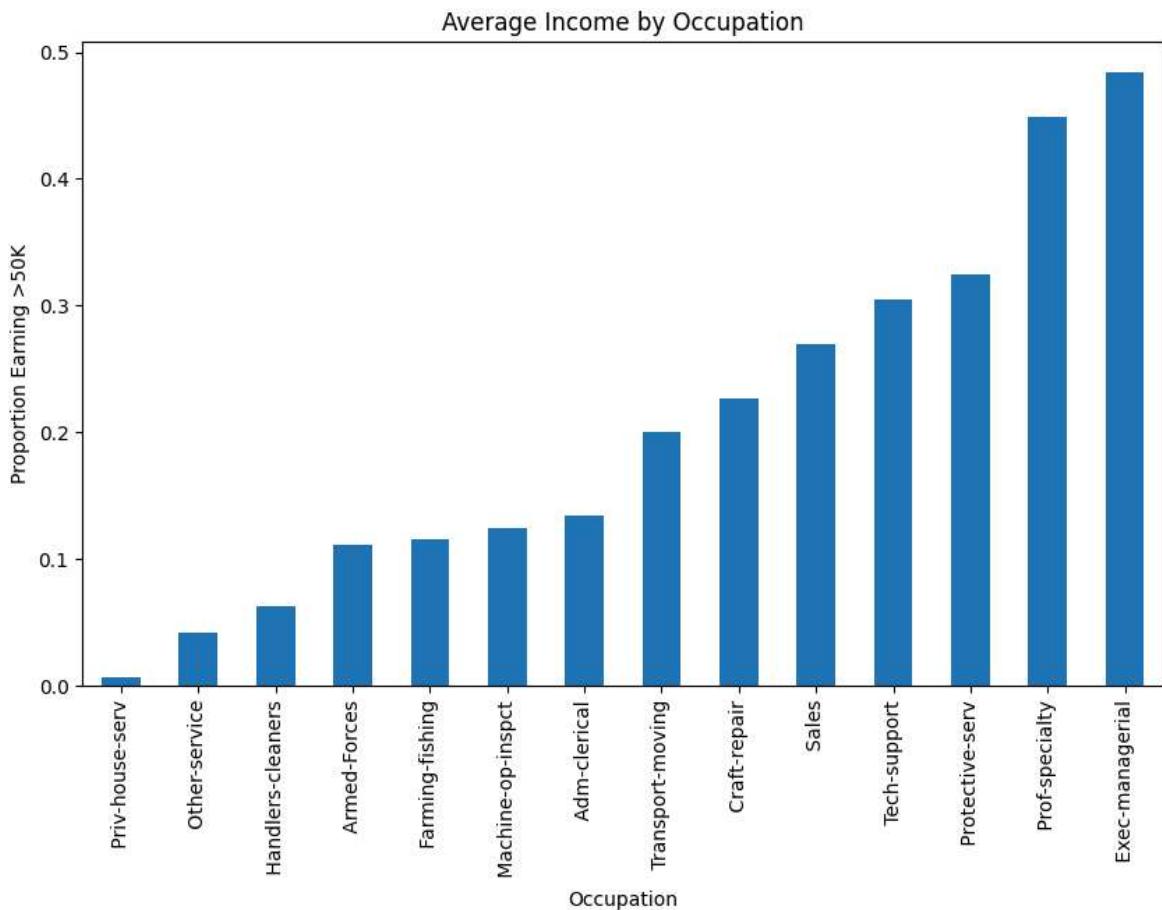
# Convert income to binary
df['income_binary'] = df['income'].apply(lambda x: 1 if x == '>50K' else 0)

# Remove missing occupation values if any
df = df[df['occupation'] != '?']

# Calculate average income by occupation
occupation_income = df.groupby('occupation')['income_binary'].mean().sort_values

# Plot bar chart
plt.figure(figsize=(10,6))
occupation_income.plot(kind='bar')
plt.title("Average Income by Occupation")
plt.ylabel("Proportion Earning >50K")
plt.xlabel("Occupation")
plt.xticks(rotation=90)
plt.show()

# Print highest and lowest
print("Highest Average Income Profession:", occupation_income.idxmax())
print("Lowest Average Income Profession:", occupation_income.idxmin())
```



Highest Average Income Profession: Exec-managerial

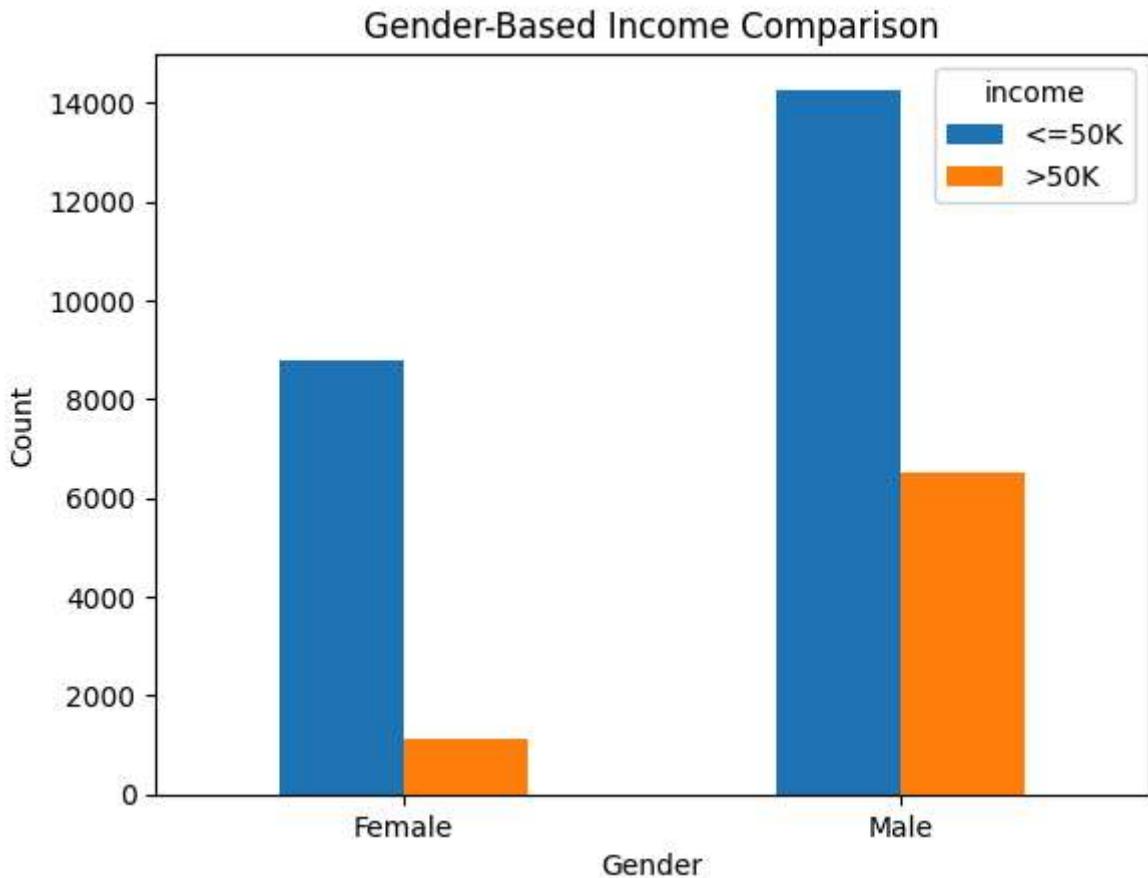
Lowest Average Income Profession: Priv-house-serv

## 7. Is there any gender-based income disparity in the dataset? Use a grouped bar chart to compare.

```
In [22]: # Clean whitespace
df.loc[:, 'income'] = df['income'].str.strip()
df.loc[:, 'gender'] = df['gender'].str.strip()

# Create grouped count table
gender_income = df.groupby(['gender', 'income']).size().unstack()

# Plot grouped bar chart
gender_income.plot(kind='bar')
plt.title("Gender-Based Income Comparison")
plt.xlabel("Gender")
plt.ylabel("Count")
plt.xticks(rotation=0)
plt.show()
```



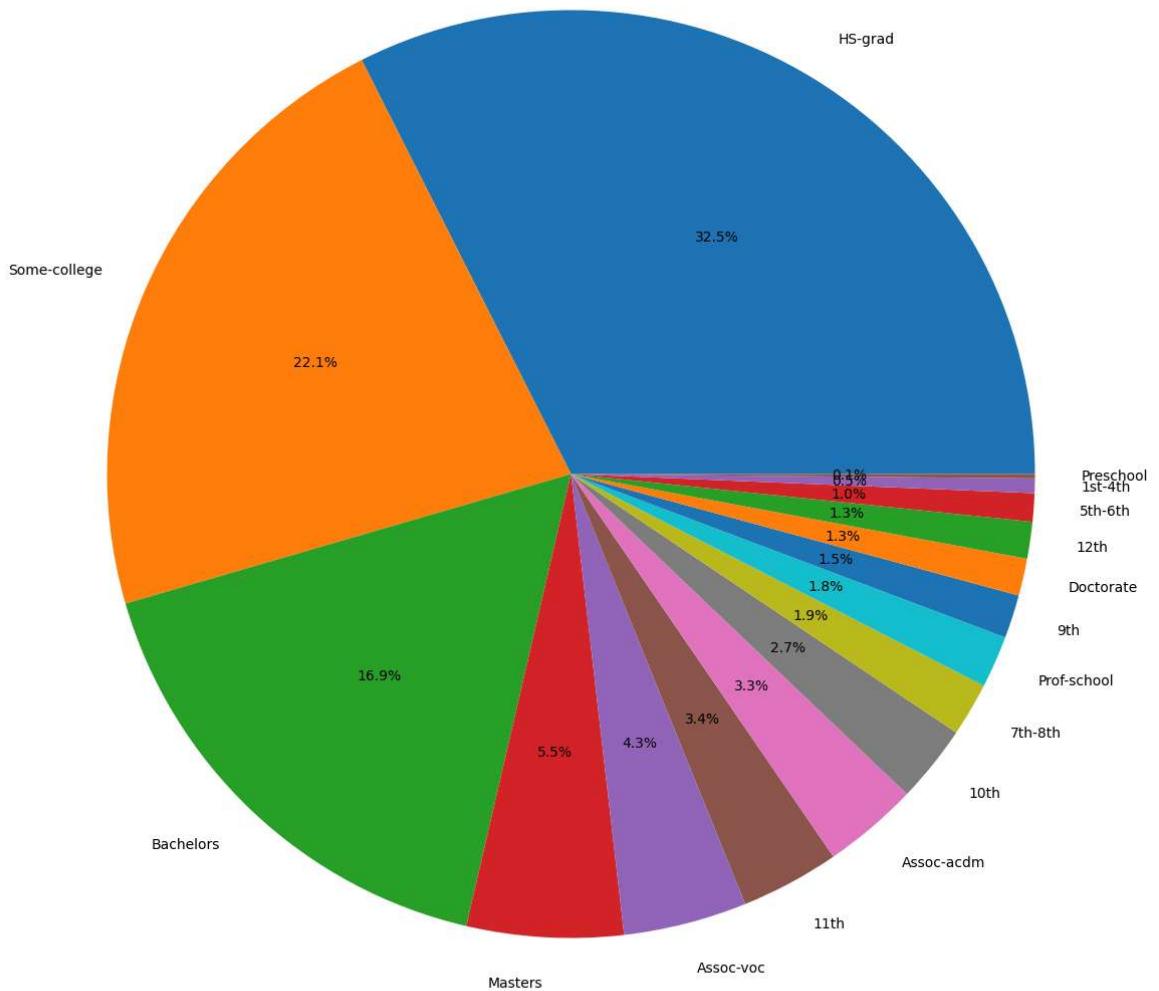
8. Can you create a pie chart showing the proportion of different education levels in the dataset?

```
In [30]: # Clean whitespace
df.loc[:, 'education'] = df['education'].str.strip()

# Count education levels
education_counts = df['education'].value_counts()

# Plot pie chart
plt.figure(figsize=(15,15))
education_counts.plot(kind='pie', autopct='%1.1f%%')
plt.title("Proportion of Education Levels")
plt.ylabel("")
plt.show()
```

Proportion of Education Levels



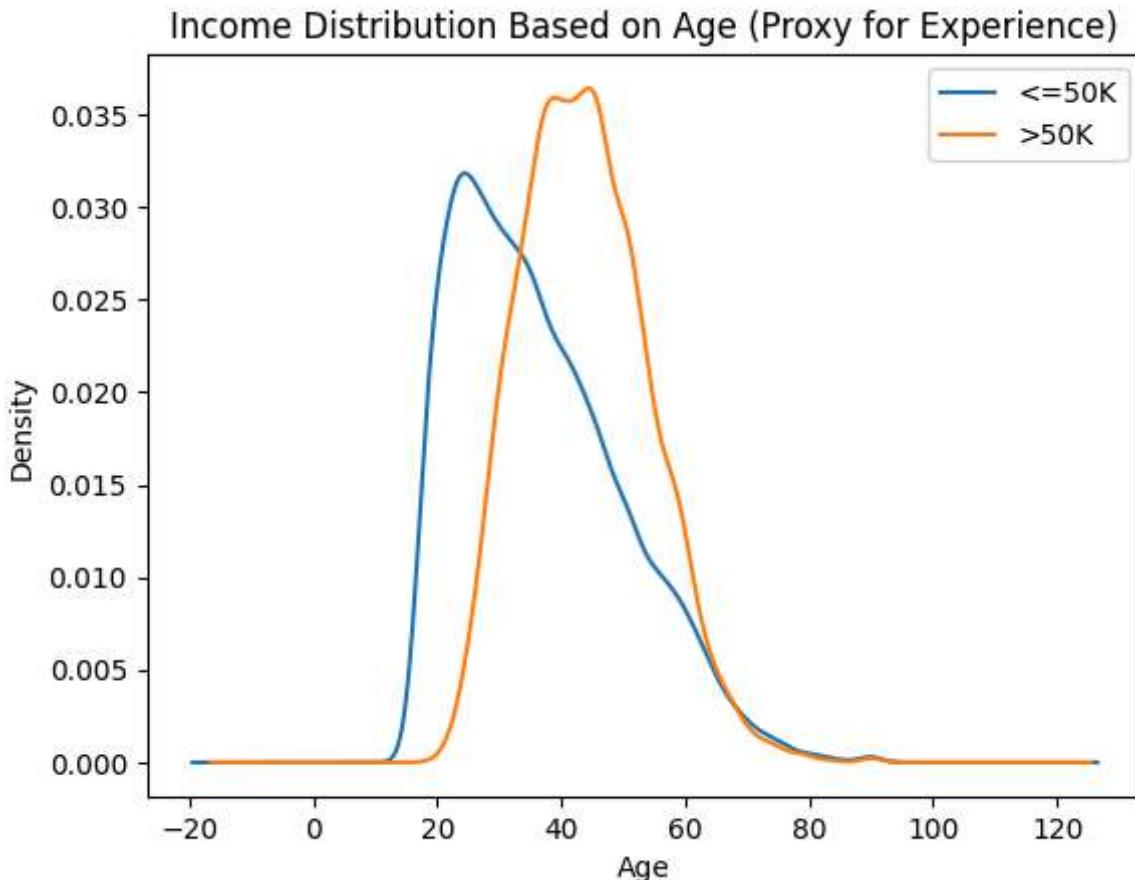
## 9. Use a violin plot or KDE plot to analyze the spread of income based on work experience.

```
In [24]: # Clean whitespace
df.loc[:, 'income'] = df['income'].str.strip()

# Separate income groups
low_income = df[df['income'] == '<=50K']['age']
high_income = df[df['income'] == '>50K']['age']

# KDE plots
low_income.plot(kind='kde')
high_income.plot(kind='kde')

plt.legend(['<=50K', '>50K'])
plt.title("Income Distribution Based on Age (Proxy for Experience)")
plt.xlabel("Age")
plt.show()
```

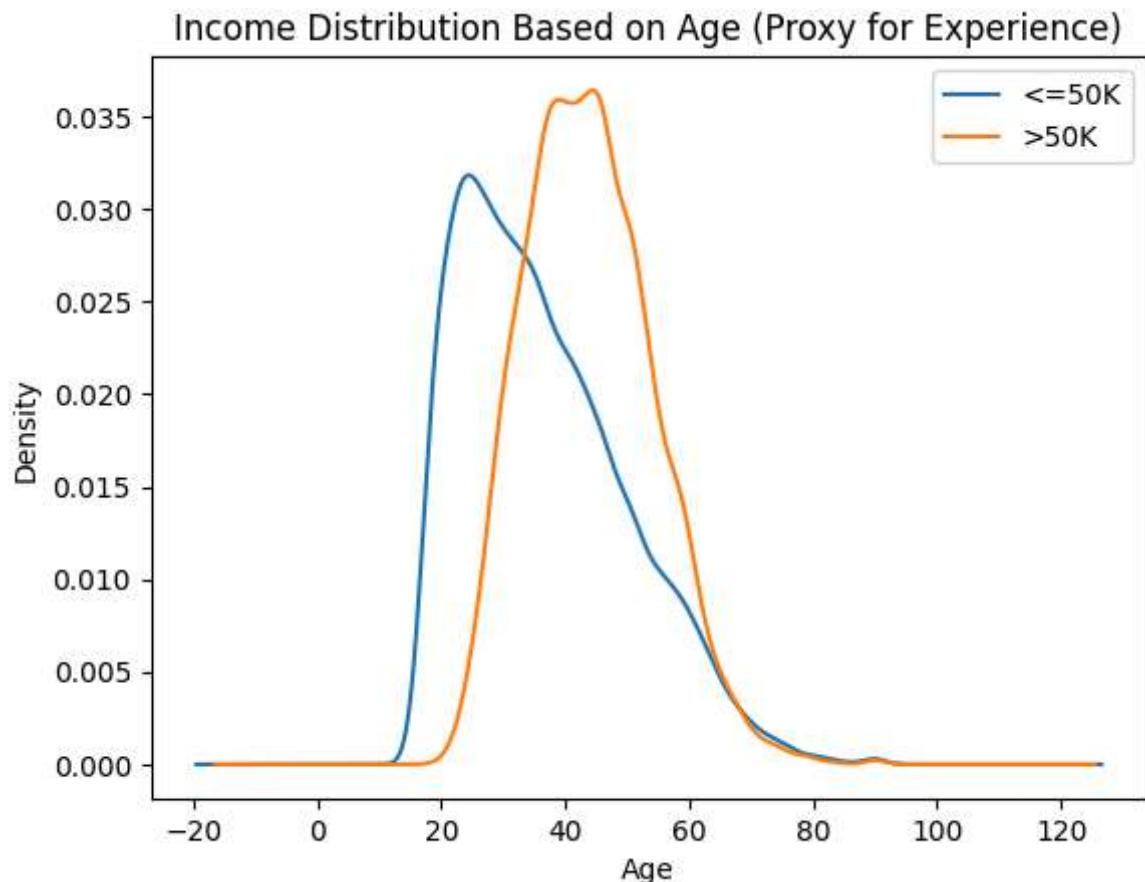


```
In [25]: # Clean whitespace
df['income'] = df['income'].str.strip()

# Separate income groups
low_income = df[df['income'] == '<=50K']['age']
high_income = df[df['income'] == '>50K']['age']

# KDE plots
low_income.plot(kind='kde')
high_income.plot(kind='kde')

plt.legend(['<=50K', '>50K'])
plt.title("Income Distribution Based on Age (Proxy for Experience)")
plt.xlabel("Age")
plt.show()
```



10. Create an interactive dashboard (using tools like Tableau or Power BI) to analyze multiple variables together.